

Lab 2 – SimpleDB Report

Shubhika Garg
Qin Zheng

TupleDesc.java

- **Time Spent:** 5 hours
- **Tasks & Solutions**
 - **TupleDesc ()**
 - An ArrayList is used because it is a dynamic array and we do not have to specify the size while creating it, the size of the array automatically increases when we dynamically add and remove data items.
 - **numFields()**
 - Used the size() function to return the number of fields in the given TupleDesc
 - **getFieldType(), fieldNameToIndex()**
 - Used the get() method to retrieve value at a specific index.
 - **getSize()**
 - We did addition of the size of tuples and stored in a variable for size of tuple in bytes.
 - **TupleDesc merge()**
 - An iterator is used to retrieve elements one by one.
 - The hasNext() method returns true if the iteration has more elements.
 - **toString()**
 - A StringBuilder is used as it creates a mutable sequence of characters. The append function appends the specified string at the end of a string.
 - **equals()**
 - Usage of instanceof keyword for checking if a reference variable is containing a given type of object reference or not. Here, it checks if object o is not an instanceof TupleDesc.
- **Problems**
 - In the method merge(), initially we initialized two loop variables with value 0 and used it in td1 and td2. We then re-read the comments which stated to have last fields of td2. We then removed the second loop variable.
 - '>=' used in some places instead of '>=' and missed curly braces in some conditions. It gave error: illegal start of expression.

Tuple.java

- **Time Spent:** 1.5 hours
- **Tasks & Solutions**
 - **toString()**
 - A StringBuilder is used as it creates a mutable sequence of characters. The append function appends the specified string at the end of a string.
 - **Iterator<Field> fields()**
 - The asList() method is used to return a fixed-size list backed by the specified array.
 - **getTupleDesc(), getRecordId(), setRecordId(), setField(), getField(), resetTupleDesc()**
 - Used this keyword to refer to the current class objects.

Catalog.java

- **Time Spent:** 3 hours
- **Tasks & Solutions**
 - **Catalog()**
 - Initialized the ArrayList variables.
 - **addTable()**
 - Used set() method to replace element present at loop variable index with method parameter values.
 - **getTableId()**
 - If the method parameter name is not null and is equal to the specified name of the list, then return the DBFile of that specific index.
 - **getTupleDesc, getDatabaseFile, getPrimaryKey() and getTableName()**
 - Used the get() method to retrieve value at a specific index.
 - **Iterator()**
 - Used an iterator() to retrieve elements one by one.
 - **clear()**
 - Initialized the ArrayList variables again.
- **Problems**
 - In the method addTable(), initially used the same loop variable for both the while loops without re-initializing the variable with 0 before using it in the second while loop.

BufferPool.java

- **Time Spent:** 1 hour
- **Tasks & Solutions**
 - **BufferPool()**
 - Defined the LinkedList and the HashMap variables and also the numPages variable.
 - **getPage()**
 - Used the getCatalog() method of the Connection interface to return the name of the current catalog /database, of the current connection object.
 - Used the addFirst() method to insert a specific pid at the beginning of a LinkedList.
 - Used put() method to associate the specified value (p) with the specified key (pid) in the map.

HeapPageId.java

- **Time Spent:** 1 hour
- **Tasks & Solutions**
 - **hashCode()**
 - We use Objects.hash to implement it
 - **equals()**
 - Check if the type, page number and table id of current pageId are the same with those of another object
- **Problems**
 - For equals(), we didn't check the instance type of the comparing object at first and it failed the test case.

RecordId.java

- **Time Spent:** 40 minutes
- **Tasks & Solutions**
 - **equals()**
 - checking if the type, tuple numbers and page id of current RecordId are equal with those of another object.
 - **hashCode()**
 - We used Objects.hash.

HeapPage.java

- **Time Spent:** 2.5 hours
- **Tasks & Solutions**
 - **getNumTuples()**
 - The number of bits used by a tuple is this.td.getSize()*8+1, the extra one bit is for indicating whether the tuple is used.
 - The total number of bits in one page can be obtained using BufferPool.getPageSize() * 8.0
 - As the tuples in a page should all be complete, we used (int) Math.floor() to round it down.
 - **getHeaderSize()**
 - The size of header is linked with the number of tuples in the page, so the number of bytes used by the header is this.getNumTuples() * 1.0 / 8.0.
 - As each tuple takes a bit in header, we use (int) Math.ceil() to round it up.
 - **isSlotUsed()**
 - The corresponding byte for the ith slot can be obtained using header [i / 8]
 - The offset in the byte for the slot can be obtained with 1 << (i % 8)
 - For example, when the slot_id is 26, then the bit in the header indicating its usage is the second bit from high to low, in the 3rd byte (start from 0)
 - **iterator()**
 - Add tuples used to the list
- **Problems**
 - For isSlotUsed(int i), at first we tried to use a for loop to check the bitmap and misunderstood the order from lower to higher, but it turned out using ">>" is much better.
 - For iterator(), we forgot to check if the slot is empty at first.

HeapFile.java

- **Time Spent:** 4 hours
- **Tasks & Solutions**
 - **getId()**
 - Using absolute file name for hashing

- **readPage(PageId pid)**
 - The method is only called by BufferPool.getPage()
 - Using RandomAccessFile.seek(), we are able to seek from the offset rather than from the beginning
 - with RandomAccessFile.read(), we can put data into buffer
 - For iterator(TransactionId tid), we create a new class named HeapFileIterator
- **readNext()**
 - First, check if it is the last tuple in the page, if yes, go to the next page
 - Then in the while loop, we check if the offset exceeds the range of current page, if yes, we go to the next page
 - If this is the last page, then we stop iterating
- **Problems**
 - For readPage(PageId pid), we misunderstood the offset in RandomAccessFile.seek() and RandomAccessFile.read() for the same thing at first and it failed.
 - For readNext(), we forgot a lot of things, yet still pass the HeapFileReadTest. Till doing ScanTest we knew how wrong we were.

SeqScan.java

- **Time Spent:** 4 hours
- **Tasks & Solutions**
 - **reset(int tableid, String tableAlias)**
 - Sets a new TupleDesc object.
- **Problems**
 - We forgot to add one when turning the page in HeapFile.readNext()
 - We forgot to check if the page number is within current page and if it is null.

